

COMPUTER ARCHITECTURE

LAB 2

Tsirakis Orestis (9995)

Kogia Iliana (10090)

Section 1

1) MinorCPU Cache

[system]

cache_line_size=64

L1_Instruction Cache: [system.cpu.icache]

assoc=2

size=32768=32kB

L1_Data Cache: [system.cpu.dcache]

assoc=2

size=65536=64kB

L2 Cache: [system.l2]

assoc=8

size=2097152=2MB

2)

Benchmarks	sim_seconds	system.cpu.cpi	system.cpu.dcache.overall_miss_rate::total	system.cpu.icache.overall_miss_rate::total	system.l2.overall_miss_rate::total
------------	-------------	----------------	--	--	------------------------------------

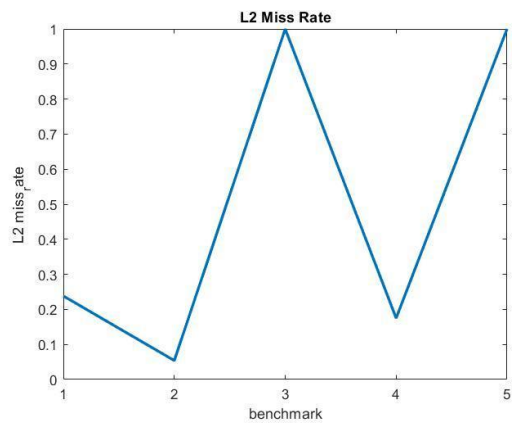
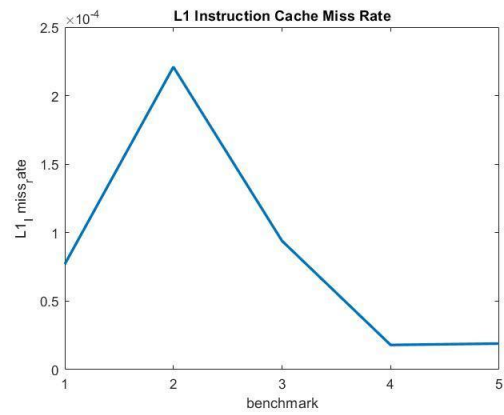
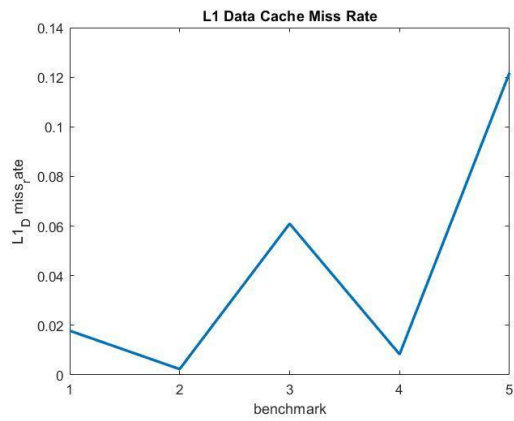
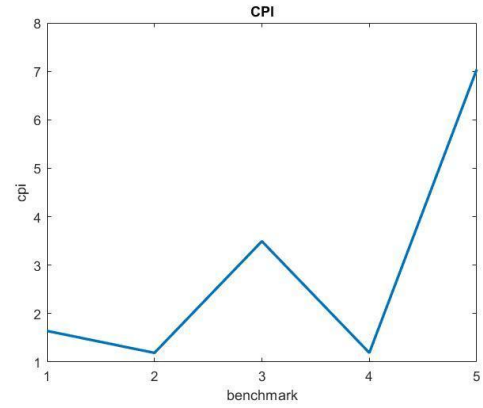
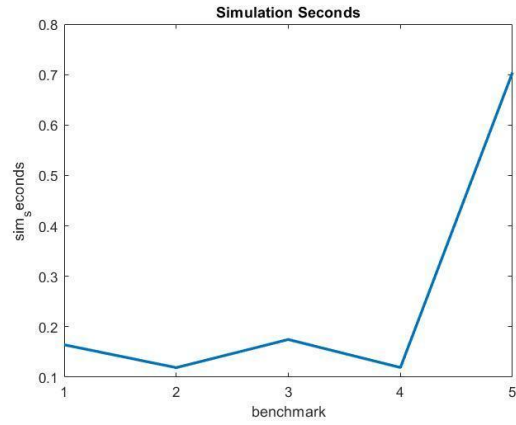
Integer benchmarks: bzip, mcf, sjeng, hmmer

Float benchmarks: libm

<https://www.spec.org/cpu2006/Docs/>

1.specbzip	0.164249	1.642495	0.017821	0.000077	0.238176
2.specbmmer	0.118876	1.188762	0.002353	0.000221	0.054059

3.speclibm	0.174671	3.493415	0.060972	0.000094	0.999944
4.specmcf	0.119056	1.190561	0.008329	0.000018	0.174927
5. specsjeng	0.704128	7.041281	0.121835	0.000019	0.999917



Benchmark no.5 (specsjeng) takes the most time to complete with 0.7 seconds simulated and has the biggest L1 Data cache miss rate.

Benchmark no.2 (spechmmer) has the most L1 instruction cache misses.

Benchmark no.3 (speclibm) has the biggest L2 miss rate.

3)

system.clk_domain.clock	1000	# Clock period in ticks
system.cpu_clk_domain.clock	500	# Clock period in ticks
1GHz:		
system.clk_domain.clock	1000	# Clock period in ticks
system.cpu_clk_domain.clock	1000	# Clock period in ticks
3GHz:		
system.clk_domain.clock	1000	# Clock period in ticks
system.cpu_clk_domain.clock	333	# Clock period in ticks

We observe that the system clock is default and includes the memory controller and memory bus, the CPU clock follows the values we give as input flags and sets the cpu core and cache clock. If we add another cpu, it will take the value of cpu cluster clock.

1GHz:

Benchmarks	sim_seconds	system.cpu.cpi	system.cpu.dcache.overall_miss_rate::total	system.cpu.icache.overall_miss_rate::total	system.l2.overall_miss_rate::total
specbzip_1GHz	0.161025	1.610247	0.014675	0.000077	0.282157
specmcf_1GHz	0.127942	1.279422	0.002108	0.023627	0.055046
spechmmer_1GHz	0.118530	1.185304	0.001629	0.000221	0.077747
specsjeng_1GHz	0.704056	7.040561	0.121831	0.000020	0.999972
speclibm_1GHz	0.262327	2.623265	0.060971	0.000094	0.999944

3GHz:

Benchmarks	sim_seconds	system.cpu.cpi	system.cpu.dcache.overall_miss_rate::total	system.cpu.icache.overall_miss_rate::total	system.l2.overall_miss_rate::total
------------	-------------	----------------	--	--	------------------------------------

We simulated the benchmarks with L1 being 32, 64, 80kB, because most of the times L1 size is close to 64kB to reduce the misses and have a higher hit rate with not much cost. We tried a bigger L1_I size than L1_D (65kB, 16kB) as the processor is usually sitting idle while an instruction is being fetched but we did not see any impact on the performance.

L2 = 256kB-4MB

We simulated the benchmarks with L2 being 256kB, 512kB, 2MB, although L2 size being 2MB may be too big for real systems.

Cache Line Size = 32, 64, 128

We simulated the benchmarks with Cache Line Size being 64B and 128B. 64B being the standard for modern systems, but we wanted to try a bigger line size to check if it improves the performance.

Associativity = 1, 2, 4

For L1 we checked associativity of 1(direct mapped) and 2 (2-way associative) as its size is smaller, for L2 we checked 2-way and 4-way associative.

2)

Benchmarks	sim_seconds	system.cpu.cpi	system.cpu.dcache.overall_miss_rate::total	system.cpu.icache.overall_miss_rate::total	system.l2.overall_miss_rate::total
------------	-------------	----------------	--	--	------------------------------------

specbzip_1	0.168500	1.684997	0.017820	0.000077	0.326983
specbzip_2	0.171925	1.719245	0.021676	0.000072	0.266583
specbzip_3	0.171944	1.719439	0.021675	0.000084	0.266600
specbzip_4	0.173792	1.737925	0.017806	0.000077	0.424388
specbzip_5	0.164249	1.642495	0.017821	0.000077	0.238176
specbzip_6	0.177353	1.773530	0.021661	0.000072	0.348441
specbzip_7	0.172442	1.724417	0.021385	0.000090	0.271747
specbzip_8	0.168359	1.683594	0.017818	0.000077	0.319777
specbzip_9	0.172283	1.722835	0.021387	0.000090	0.265389
specbzip_10	0.168312	1.683117	0.017841	0.000067	0.239170

spechmmer_1	0.118877	1.188768	0.002353	0.000221	0.054355
spechmmer_2	0.119081	1.190813	0.002661	0.000087	0.048894
spechmmer_3	0.119172	1.191716	0.002661	0.000585	0.045956
spechmmer_4	0.118930	1.189300	0.002358	0.000222	0.064771

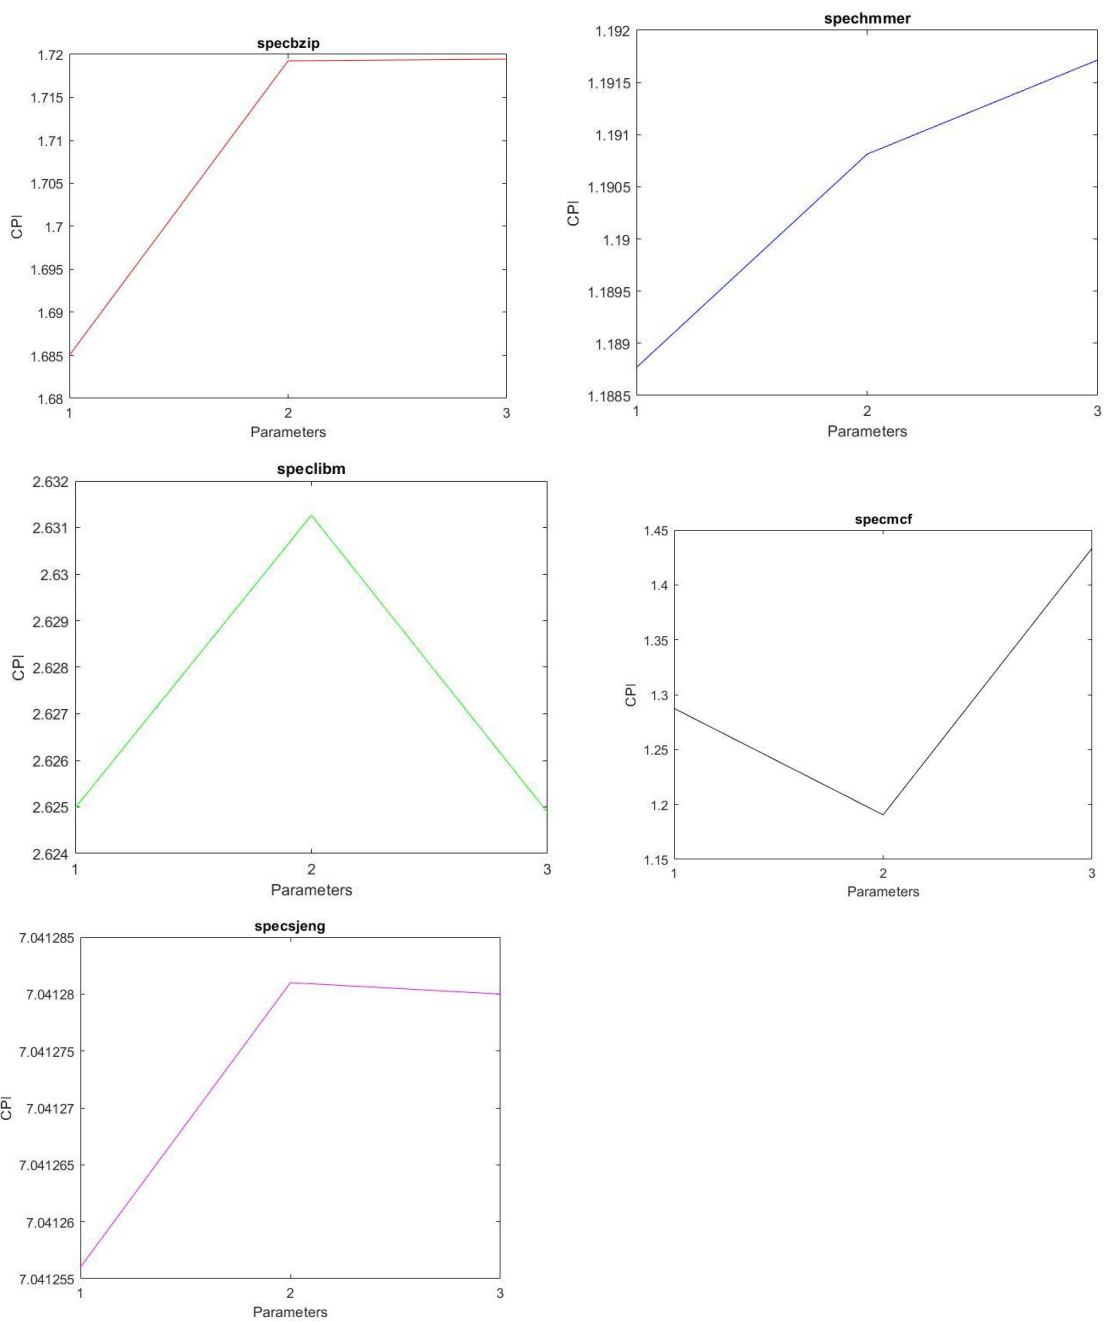
spechmmer_5	0.118876	1.188762	0.002353	0.000221	0.054059
spechmmer_6	0.119121	1.191210	0.002663	0.000087	0.054738
spechmmer_7	0.136136	1.361358	0.005676	0.034155	0.006823
spechmmer_8	0.118876	1.188762	0.002353	0.000221	0.054059
spechmmer_9	0.136132	1.361321	0.005676	0.034155	0.006728
spechmmer_10	0.118343	1.183426	0.001387	0.000347	0.048683

speclibm_1	0.262498	2.624978	0.060972	0.000094	0.999942
speclibm_2	0.263126	2.631263	0.060972	0.000088	0.999959
speclibm_3	0.262488	2.624883	0.060972	0.000113	0.999858
speclibm_4	0.262529	2.625290	0.060972	0.000094	0.999942
speclibm_5	0.262319	2.623187	0.060972	0.000094	0.999942
speclibm_6	0.262547	2.625472	0.060972	0.000088	0.999959
speclibm_7	0.265035	2.650347	0.062151	0.000122	0.973222
speclibm_8	0.262498	2.624978	0.060972	0.000094	0.999942
speclibm_9	0.265002	2.650021	0.062151	0.000122	0.973222
speclibm_10	0.199147	1.991471	0.030487	0.000112	0.999790

specmcf_1	0.128751	1.287506	0.002493	0.023626	0.063724
specmcf_2	0.119056	1.190561	0.008329	0.000018	0.174927
specmcf_3	0.143351	1.433507	0.008332	0.042807	0.031770
specmcf_4	0.129013	1.290132	0.002493	0.023626	0.067543
specmcf_5	0.128163	1.281627	0.002493	0.023626	0.054518
specmcf_6	0.119280	1.192799	0.008329	0.000018	0.184985
specmcf_7	0.134142	1.341420	0.006987	0.034838	0.039084
specmcf_8	0.128706	1.287057	0.002493	0.023626	0.063051
specmcf_9	0.134093	1.340926	0.006987	0.034838	0.038598
specmcf_10	0.132824	1.328239	0.001859	0.034828	0.027021

specsjang_1	0.704126	7.041256	0.121833	0.000020	0.999950
specsjang_2	0.704128	7.041281	0.121835	0.000019	0.999917

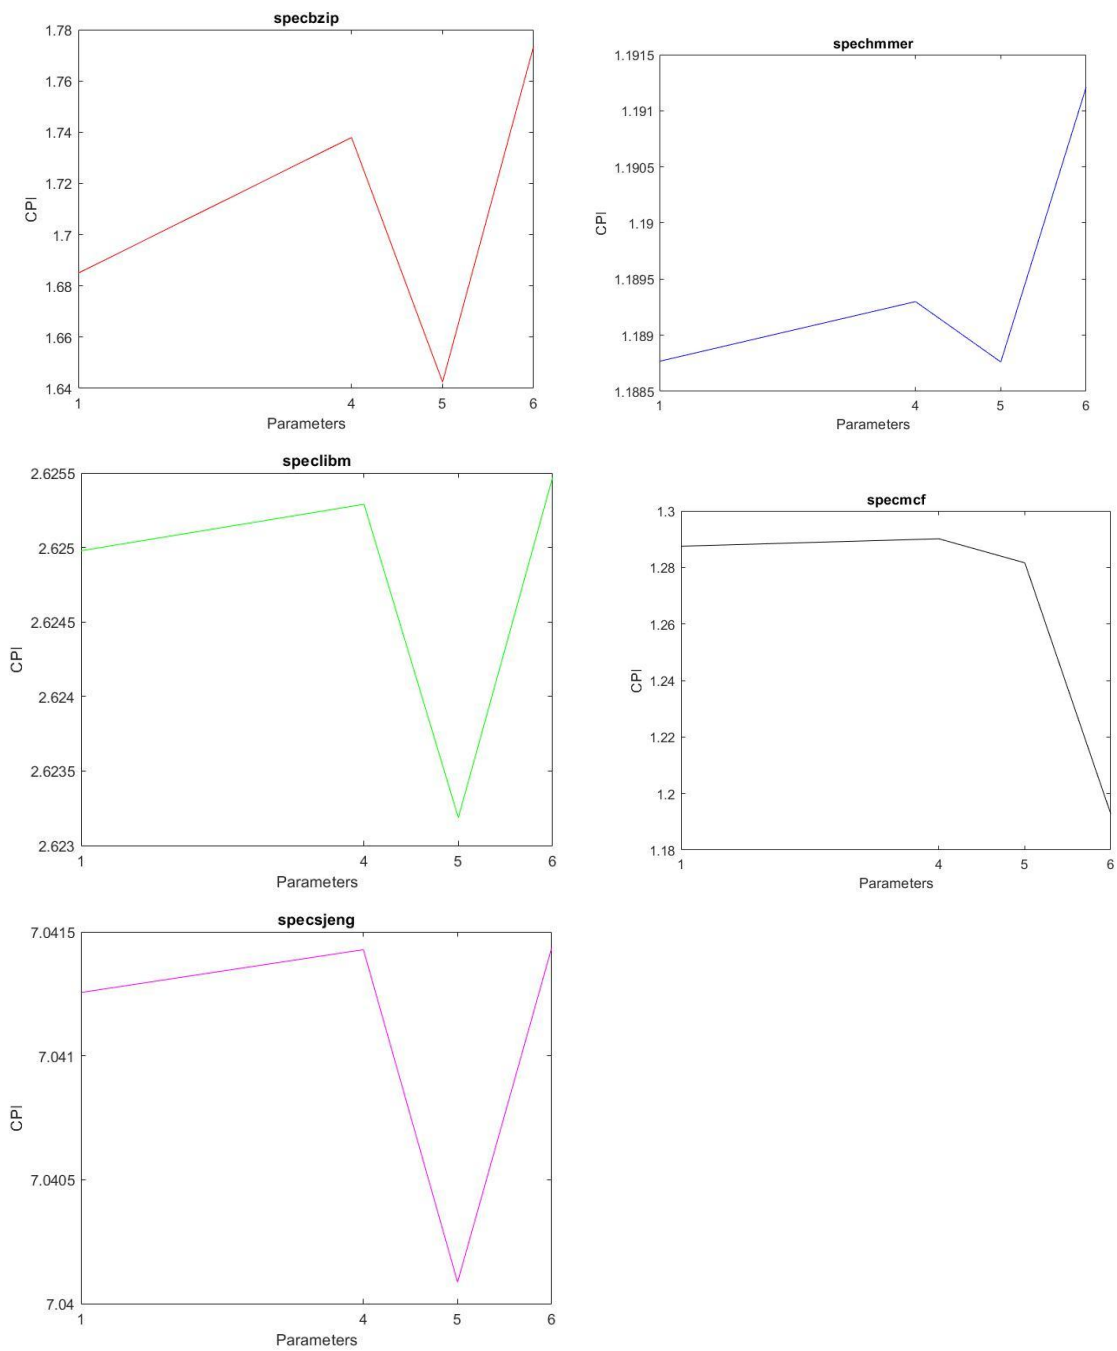
{1, 2, 3}



We observe that for the L1 sizes of the set no.1 we get the optimum CPI for all benchmarks except specmcf (for mcf the best is set no.2).

L1_D=32kB, L1_I=32kB

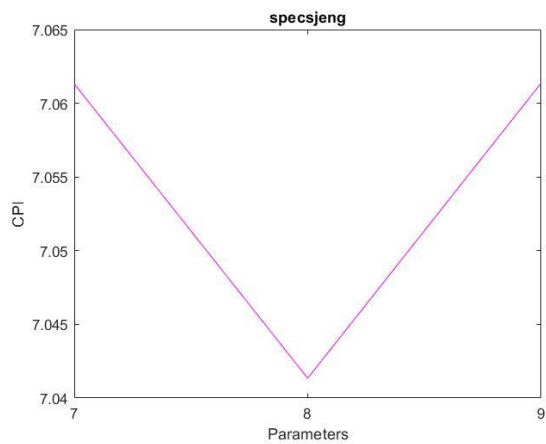
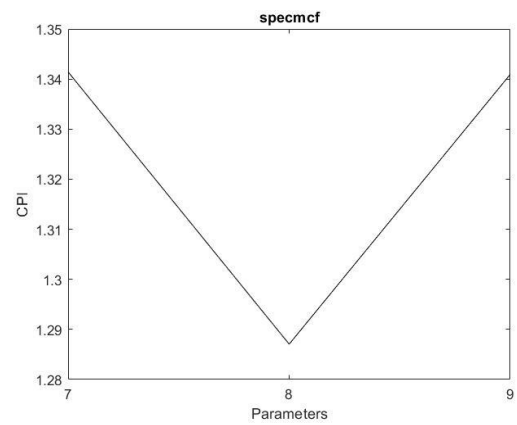
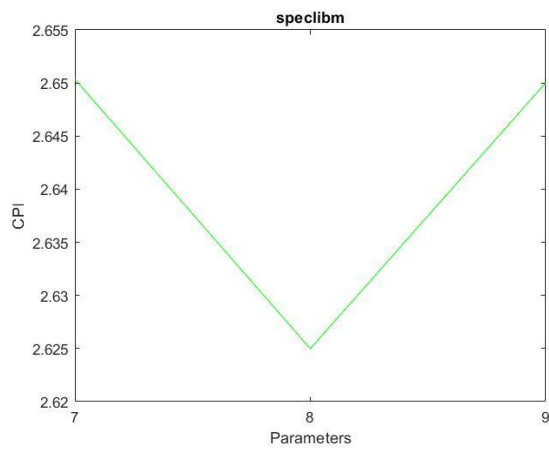

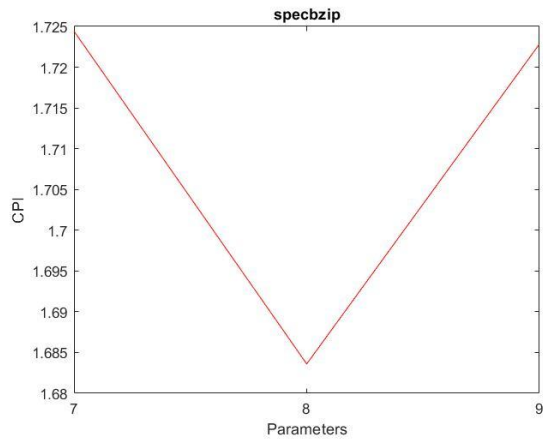
{1, 4, 5, 6}



We observe that for the L2 sizes of the set no.5 we get the optimum CPI for all benchmarks except specmcf.

L2=2MB.

{7, 8, 9}



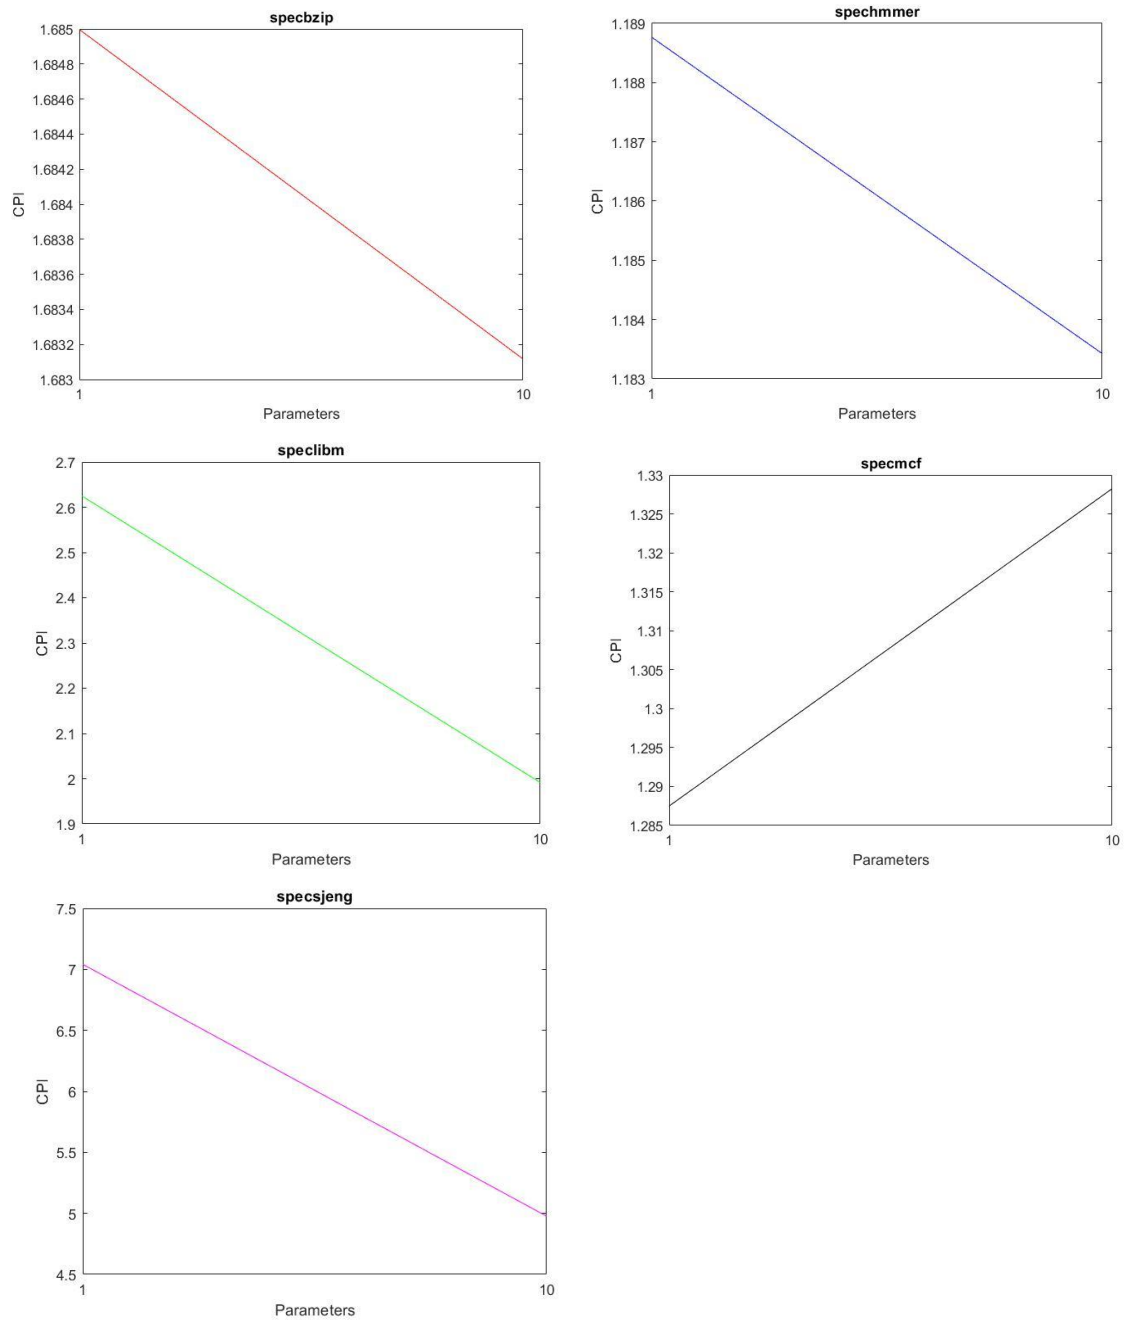
We observe that for the L1_Ass and L2_Ass of the set no.8 we get the optimum CPI for all benchmarks.

L1_D_Ass=2

L1_I_Ass=2

L2_Ass=4

{1, 10}



We observe that for the Cache line sizes of the set no.10 we get the optimum CPI for all benchmarks except specmcf.

Cache_Line_Size = 128

For benchmark MCF we mostly get the best performance for different values but we chose the optimal values for all the benchmarks.

Max Performance:

L1_Data = 32kB

L1_Instruction = 32kB

L2 = 2MB

L1_D_Ass = 2

L1_I_Ass = 2

L2_Ass = 4

Cache_line_size = 128

Benchmarks	sim_seconds	system.cpu.cpi	system.cpu.dcache.overall_miss_rate::total	system.cpu.icache.overall_miss_rate::total	system.l2.overall_miss_rate::total
speckzip_best	0.163596	1.635955	0.017843	0.000067	0.138201
spechmmer_best	0.118341	1.183408	0.001387	0.000347	0.048161
speclibm_best	0.199046	1.990458	0.030487	0.000112	0.999790
specmcf_best	0.132229	1.322294	0.001858	0.034834	0.020409
specsjang_best	0.497481	4.974806	0.060926	0.000015	0.999685

We get the best performance from all the other simulations for each benchmark.

Section 3

c1 = cost unit.

$$\text{Cost} = c1 * L1_size + (c1/10^2) * L2_size + 2 * c1 * L1_ass/16 + c1 * L2_ass/16 + c1 * \text{Cache_L_S}/8$$

The cost of L1 must be larger than the cost of L2.

As the size of the memory chip increases the cost increases too.

Increasing the associativity has a large effect on the complexity of the memory unit, so the cost rises.

The Cache Line Size costs a lot per byte as the more words saved in the same line the larger the complexity for the multiplexer to find the right word.

a) L1 = 64kB, L2 = 512kB, Ass = 2,2,4

b) L1 = 64kB, L2 = 2MB, Ass = 2,2,4

- a) $\text{Cost a} = 64 * c_1 + 5,12 * c_1 + 0.25 * c_1 + 0.25 * c_1 + 8 * c_1 = 77.62 * c_1$
- b) $\text{Cost b} = 64 * c_1 + 20 * c_1 + 0.5 * c_1 + 16 * c_1 = 100 * c_1$

The small difference in efficiency of L2 = 2MB and Cache_Line_size = 128 is not worth such a bigger cost we will have to pay to implement our system. So we think it will be best to consider an L2 = 512kB and Cache_Line_Size of 64 to allow both efficiency and price availability.